

**High Performance Computing**  
**Anthony Trubiano**  
**Assignment 3 (due April 1, 2019)**

**Note:** All computations were performed with an Intel Core i7-8750H Processor with base frequency 2.20 GHz. The maximum main memory bandwidth is 41.8 GB/s. At 16 double precision operations per cycle, the theoretical max flop rate would be about 35.2 GFlops/s. It has 6 cores and can reach 12 threads through hyper-threading.

1. **Approximating Special Functions Using Taylor Series and Vectorization.** I improved the accuracy of the function `sin4_vec()` simply by adding more terms to the Taylor series using the `Vector` class. Including terms up to  $O(x^{11})$  gives the same error as the `sin4_taylor()` code given but runs about 2.5 times faster.
2. **Parallel Scan in OpenMP** Here we parallelize the scan operation using  $p$  threads. Let  $n$  be the length of the vector. The parallel implementation takes in the number of threads, then assigns each a chunk of size  $p/n$ . If there is a remainder, the last threads gets assigned these entries. Each thread then performs the scan operation on its chunk. We then compute a correction in serial by adding the final entry of each chunk to all entries in the following chunk. The speed of the algorithm as a function of number of threads is shown in the following table for  $n = 100,000,000$ .

Threads	Time (s)
Serial	0.224
1	0.287
2	0.188
3	0.153
4	0.141
5	0.135
6	0.132

**Table 1:** Time taken (in seconds) for the scan operation on a vector of length  $n = 100,000,000$  as a function of thread number.

From this table we see that using 2 threads is enough to get slightly better performance while using 4 threads approximately halves the amount of time taken. We don't see  $p$  times speed up because the parallel implementation has to perform the fix calculations that the serial version does not, which takes a non-negligible amount of time. Additionally, if we make  $n$  smaller, by a factor of 10 for example, the parallel version is about the same if not slower than the serial version. We need a large enough vector to get noticeable improvements from parallelization.